

## HOMEWORK: HOMEWORK 4.2

Suppose you have a collection called *tweets* whose documents contain information about the *created\_at* time of the tweet and the user's *followers\_count* at the time they issued the tweet. What can you infer from the following *explain* output?

```
db.tweets.find({"user.followers_count":{"$gt:1000}}).sort({"created_at" : 1 }).limit(10).skip(5000).explain()
{
  "cursor" : "BtreeCursor created_at_-1 reverse",
  "isMultiKey" : false,
  "n" : 10,
  "nscannedObjects" : 46462,
  "nscanned" : 46462,
  "nscannedObjectsAllPlans" : 49763,
  "nscannedAllPlans" : 49763,
  "scanAndOrder" : false,
  "indexOnly" : false,
  "nYields" : 0,
  "nChunkSkips" : 0,
  "millis" : 205,
  "indexBounds" : {
    "created_at" : [
      [
        {
          "$minElement" : 1
        },
        {
          "$maxElement" : 1
        }
      ]
    ]
  },
  "server" : "localhost.localdomain:27017"
}
```



- ☒ This query performs a collection scan.
- ☒ The query uses an index to determine the order in which to return result documents.
- ☐ The query uses an index to determine which documents match.
- ☐ The query returns 46462 documents.
- ☒ The query visits 46462 documents.
- ☐ The query is a "covered index query".

[SHOW ANSWER](#)

*Submission deadline has passed.*

